

Commands explanation setup

This is a copy of a conversation between ChatGPT & Anonymous.



What commands i need to run
How i do
Explain

Got it 👍 Let's break this down step by step so you can quickly set up your **Jupyter Notebook (.ipynb)** for the Internal Practical Exam. I'll give you the **commands** and explain **what to put in each section**.



Structure You Need in the Notebook

From the instructions in the screenshot:

1. **Print library versions**
2. **Reproducibility block** (install commands, seeds, CPU/GPU info, timestamp, student name & roll no)
3. **Model training & evaluation** (accuracy vs fairness)
4. **Bias findings & fairness metrics**
5. **Bias mitigation** (Fairlearn / balancing / feature removal)
6. **Ethical implications & short conclusion** (<300 words)
7. **All tables, metrics, confusion matrices, plots must be generated inside the notebook**

♦ Step 1: Import Libraries & Print Versions

python

```
import numpy as np
import pandas as pd
import sklearn
import torch
import transformers
import nltk
import fairlearn

print("numpy:", np.__version__)
print("pandas:", pd.__version__)
print("scikit-learn:", sklearn.__version__)
print("torch:", torch.__version__)
print("transformers:", transformers.__version__)
print("nltk:", nltk.__version__)
print("fairlearn:", fairlearn.__version__)
```

◆ Step 2: Reproducibility Block

```
python

import random, os, time
import numpy as np
import torch

# Fix random seeds
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)

# Check device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

# Student Info + Timestamp
print("Student Name: YOUR_NAME")
print("Roll Number: YOUR_ROLL_NO")
print("Timestamp:", time.strftime("%Y-%m-%d %H:%M:%S", time.localtime()))
```

◆ Step 3: Load Dataset (Example: sklearn dataset)

If your teacher didn't specify dataset, you can use something like **Adult Income dataset** (commonly used for fairness testing).

```
python

from sklearn.datasets import fetch_openml

# Fetch Adult Income dataset (for fairness demo)
data = fetch_openml(data_id=1590, as_frame=True) # Adult dataset
df = data.frame
df.head()
```

◆ Step 4: Preprocess Data & Train Model

```
python

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Target: income
X = df.drop(columns=["class"])
y = df["class"]

# Encode categorical columns
X = X.apply(lambda col: LabelEncoder().fit_transform(col))

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=SEED)

# Train simple model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

# Predictions
```

```
y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

◆ Step 5: Fairness Analysis (Using fairlearn)

```
python

from fairlearn.metrics import MetricFrame, selection_rate, demographic_parity_difference

sensitive_feature = X_test["sex"] # fairness check across gender

mf = MetricFrame(metrics=accuracy_score,
                  y_true=y_test,
                  y_pred=y_pred,
                  sensitive_features=sensitive_feature)

print("Accuracy by group:", mf.by_group)
print("Overall accuracy:", mf.overall)
print("Demographic parity difference:", demographic_parity_difference(y_test, y_pred,
                             sensitive_features=sensitive_feature))
```

◆ Step 6: Bias Mitigation

Example: reweighting or removing sensitive feature

```
python

# Remove sensitive feature 'sex' and retrain
X_train_fair = X_train.drop(columns=["sex"])
X_test_fair = X_test.drop(columns=["sex"])

model_fair = LogisticRegression(max_iter=1000)
model_fair.fit(X_train_fair, y_train)
y_pred_fair = model_fair.predict(X_test_fair)

print("New Accuracy:", accuracy_score(y_test, y_pred_fair))
print("New Confusion Matrix:\n", confusion_matrix(y_test, y_pred_fair))

mf_fair = MetricFrame(metrics=accuracy_score,
                      y_true=y_test,
                      y_pred=y_pred_fair,
                      sensitive_features=sensitive_feature)

print("Fairness after mitigation (Accuracy by group):", mf_fair.by_group)
```

◆ Step 7: Plots

```
python

import matplotlib.pyplot as plt

mf.by_group.plot(kind="bar")
plt.title("Accuracy by Gender before Mitigation")
plt.show()

mf_fair.by_group.plot(kind="bar", color="green")
```

```
plt.title("Accuracy by Gender after Mitigation")  
plt.show()
```

◆ Step 8: Conclusion (≤300 words)

At the end, add a **Markdown cell**:

```
pgsql
```

```
### Conclusion  
The experiment compared accuracy and fairness of a logistic regression model on the Adult Income dataset.  
Initial results showed bias: male group had higher accuracy and selection rates compared to female group.  
Bias mitigation by removing the sensitive attribute (`sex`) reduced bias difference but slightly impacted accuracy.  
This shows a tradeoff between accuracy and fairness.  
Ethically, it highlights the need to balance performance with fairness to avoid discrimination in real-world AI applications.
```

✓ Now, when you **run notebook top-to-bottom**, you'll have:

- Library versions printed
- Reproducibility ensured
- Accuracy vs fairness comparison
- Bias mitigation demonstration
- Confusion matrices, plots, metrics
- Short conclusion

Do you want me to **make you a ready-to-run Colab notebook** with this code (so you just copy & paste), or do you prefer I give you only the commands and you build it step by step?